

# Scan

Parallel Algorithm Design WS21/22

N. Kochendörfer, C. Alles, S. Proß

February 17, 2022

# Table of Contents

- 1 Scan Theory
- 2 Implementation
- 3 Optimizations
- 4 Summary


# Scan Theory

- Synonyms: prefix sum, cumulative sum or scan
- inclusive and exclusive version
- further specialization: segmented scan

# Inclusive Scan

Prefix Sum  
P of A

→ Store  
→ Sum



<https://williamjrjibeiro.com/?p=132>

# Inclusive vs. Exclusive scan

X	3	4	6	3	8	7	5	4
---	---	---	---	---	---	---	---	---

Y	0	3	7	13	16	24	31	36
---	---	---	---	----	----	----	----	----

Exclusive Scan

Y	3	7	13	16	24	31	36	40
---	---	---	----	----	----	----	----	----

Inclusive Scan

<https://livebook.manning.com/book/parallel-and-high-performance-computing/chapter-5/v-11/>

# Segmented Variant

1	2	3	4	5	6	input
1	0	0	1	0	1	flag bits
1	3	6	4	9	6	segmented scan +

[https://en.wikipedia.org/wiki/Segmented\\_scan](https://en.wikipedia.org/wiki/Segmented_scan)

# Implementation

## STL Algorithm

STL provides:

- `std::inclusive_scan`
- `std::exclusive_scan`

Essentially equivalent to:

```
float sum = 0;
for (size_t i = 0; i < N; i++)
{
    sum += input[i];
    output[i] = sum;
}
```

⇒ Sequential to a fault!

# Implementation

## Alternatives

### Alternatives to STL:

- OpenMP: scan pragma
- TBB: parallel\_scan function

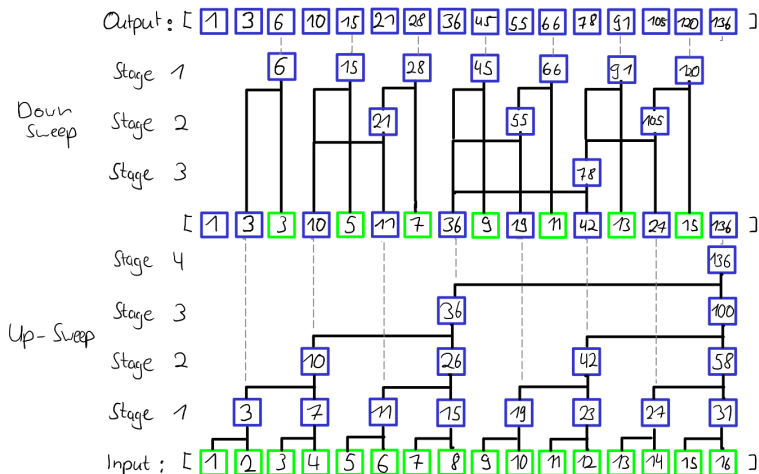
### Alternative Algorithms:

- Up-Down Sweeping Scan
- Tiled Scan



# Up-Down Sweep

## Schema Inclusive



# Up-Down Sweep

Dependency:

- Only between stages
- ⇒ Lots of parallelism

Downsides:

- Workload of  $2N$
- Communication!
- Workload stage dependent!

# Tiled Scan

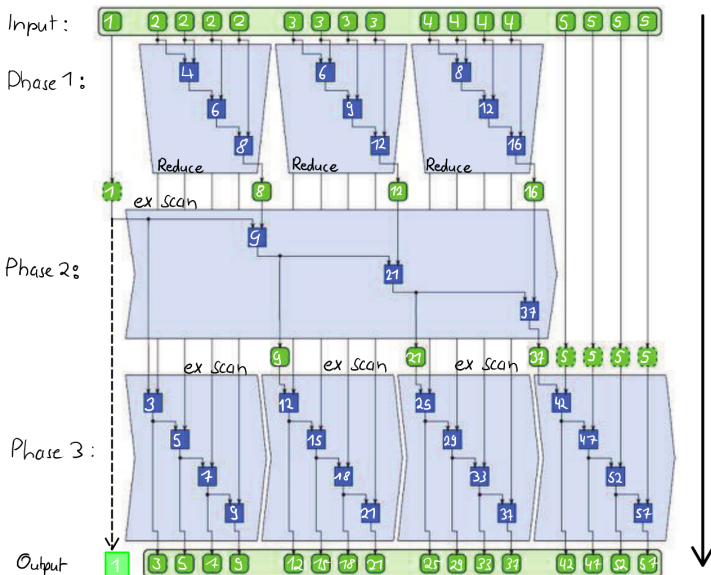
Idea: Process input in independent chunks.

- Each chunk misses previous results  
⇒ Second pass over data.
- Workload:  $2N$


Our solution:

- Temporary vector for intermediate sums.
- Only one write to output.

# Tiled Scan Schema



# Sequential Scan Results



# Segmented Scan

## Implementation

- Not present in STL!
- No reference implementations...

Solution: Wrapping the binary operation!

```
[binary_op](PairType left , PairType right){  
    PairType new_right = right;  
    if (not right.flag)  
        new_right.value =  
            binary_op(left.value , right.value);  
    return new_right;  
});
```

# Segmented Scan

## Solution


Works for:

- STL Scans
- Most inclusive scans

Challenge: Exclusive Scan


- Exclusive Segmented is complex
- Custom solution for each variant

# Sequential Segmented Scan Results






# Parallel Scan Results



# Parallel Segmented Scan Results



# Intermediate Results

Ranking:

- 1 Library provided implementations
- 2 Tiled Scan
- 3 Up-Down Sweeping Scan

Remarks:

- OpenMP  $\geq$  TBB (if available)
- Up-Down Sweep is slow

Can we do better?

# Algorithmic Optimization

Initial Goal: functional correctness.

Algorithmic Optimizations:

- Loop-Fusion:
  - Up-Down Sweep
  - Exclusive Segmented Scan
- Limiting Memory Accesses
- General clean up

Performance gain  $\sim$  1-5 GB/s!

# Data Locality

Ensure that the data generated is local to the node:

```
std::vector<float> data(N);  
#pragma omp parallel for schedule(static)  
for (size_t i = 0; i < data.size(); i++)  
{  
    data[i] = rand();  
}
```

- The performance gain by using data local structures is likely to be small due to the warmup of Catch2

# OpenMP Scheduling





Fig. 1: Different Scheduling Strategies for 100 Iterations and 4 Threads


# OpenMP Scheduling - Results MP-Media

mp-media OpenMP Provided Inclusive Scan with Different Scheduling (gnu)



# OpenMP Scheduling - Results Ziti-Rome

ziti-rome OpenMP Provided Inclusive Scan with Different Scheduling (gnu)





# TBB Partitioning


## **TBB parallel constructs used:**

- `parallel_scan`
- `parallel_for`


## **available partitioners:**

- `auto_partitioner`
- `affinity_partitioner`
- `simple_partitioner`
- `static_partitioner`

# Performance (inclusive scan)



# Performance (inclusive scan)



# Vectorization


## Requirements:

- No loop carried dependency
- Loop bounds
- No jumps in code


## Realising it:

- `#pragma omp simd`
- Compiling with `-O3`
- Using Intel `lxc` Compiler

# Vectorization - Results MP-Media



# Vectorization - Results Ziti-Rome



# Summary

Library Provided Scans are fastest

# Summary

Library Provided Scans are fastest

Optimization done:

- Algorithmic
- Data Locality
- Scheduler & Partitioner
- Vectorization


We have

- 4 versions of Scan
- 3 different algorithms
- 2 parallelization libraries + sequential

⇒ 36 Versions to keep track of!



# Roofline Mp-Media



# Roofline Ziti-Rome

