

# Scan

Parallel Algorithm Design WS21/22

N. Kochendörfer, C. Alles, S. Proß

February 16, 2022

# Table of Contents

- 1 Scan Theory
- 2 Implementation
- 3 Optimizations
- 4 Summary

2022-02-16

Scan

└─ Table of Contents

Table of Contents

- 1 Scan Theory
- 2 Implementation
- 3 Optimizations
- 4 Summary

# Scan Theory

- Synonyms: prefix sum, cumulative sum or scan
- inclusive and exclusive version
- further specialization: segmented scan

2022-02-16

Scan  
└ Scan Theory  
    └ Scan Theory


Scan Theory

- Synonyms: prefix sum, cumulative sum or scan
- inclusive and exclusive version
- further specialization: segmented scan

## Inclusive Scan

Prefix Sum  
P of A

Store  
Sum


<https://williamjrjibeiro.com/?p=132>


2022-02-16

Scan

└ Scan Theory

└ Inclusive Scan

Inclusive Scan



## inclusive vs. exclusive scan

X 

3	4	6	3	8	7	5	4
---	---	---	---	---	---	---	---

Y 

0	3	7	13	16	24	31	36
---	---	---	----	----	----	----	----

 Exclusive Scan

Y 

3	7	13	16	24	31	36	40
---	---	----	----	----	----	----	----

 Inclusive Scan

<https://livebook.manning.com/book/parallel-and-high-performance-computing/chapter-5/v-11/>

2022-02-16

Scan  
└ Scan Theory

└ inclusive vs. exclusive scan

inclusive vs. exclusive scan

X 

3	4	6	3	8	7	5	4
---	---	---	---	---	---	---	---

Y 

0	3	7	13	16	24	31	36
---	---	---	----	----	----	----	----

 Exclusive Scan

Y 

3	7	13	16	24	31	36	40
---	---	----	----	----	----	----	----

 Inclusive Scan

<https://livebook.manning.com/book/parallel-and-high-performance-computing/chapter-5/v-11/>

## segmented variant

1	2	3	4	5	6	input
1	0	0	1	0	1	flag bits
1	3	6	4	9	6	segmented scan +

[https://en.wikipedia.org/wiki/Segmented\\_scan](https://en.wikipedia.org/wiki/Segmented_scan)

2022-02-16

Scan

└ Scan Theory

└ segmented variant

segmented variant

1	2	3	4	5	6	input
1	0	0	1	0	1	flag bits
1	3	6	4	9	6	segmented scan +

[https://en.wikipedia.org/wiki/Segmented\\_scan](https://en.wikipedia.org/wiki/Segmented_scan)

# Implementation

## STL Algorithm

STL provides:

- `std::inclusive_scan`
- `std::exclusive_scan`

Essentially equivalent to:

```
float sum = 0;
for (size_t i = 0; i < N; i++)
{
    sum += input[i];
    output[i] = sum;
}
```

⇒ Sequential to a fault!

2022-02-16

Scan

└ Implementation

└ Implementation

Implementation

STL Algorithm

STL provides:

- `std::inclusive_scan`
- `std::exclusive_scan`

Essentially equivalent to:

```
float sum = 0;
for (size_t i = 0; i < N; i++)
{
    sum += input[i];
    output[i] = sum;
}
```

⇒ Sequential to a fault!

# Implementation

## Alternatives

### Alternatives to STL:

- OpenMP: scan pragma
- TBB: parallel\_scan function

### Alternative Algorithms:

- Up-Down Sweeping Scan
- Tiled Scan

2022-02-16

Scan

└ Implementation

└ Implementation

Implementation  
Alternatives

Alternatives to STL:

- OpenMP: scan pragma
- TBB: parallel\_scan function

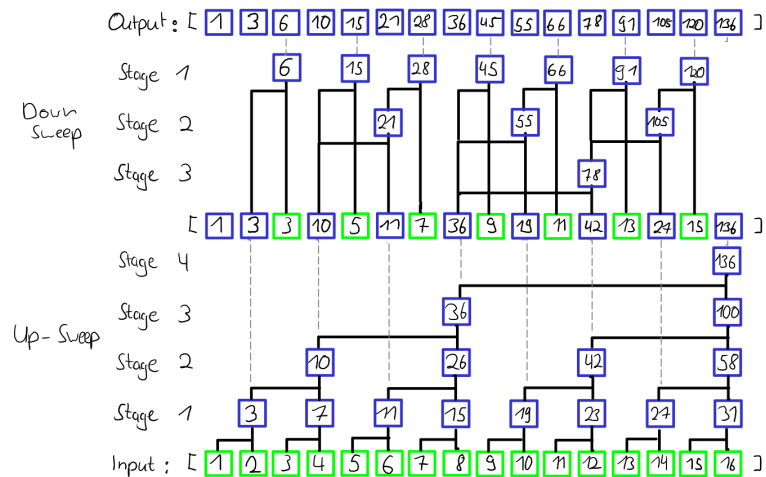
Alternative Algorithms:

- Up-Down Sweeping Scan
- Tiled Scan



## Up-Down Sweep

Schema Inclusive




2022-02-16

Scan

Implementation

Up-Down Sweep

Up-Down Sweep  
Schema Inclusive

# Up-Down Sweep

Dependency:

- Only between stages

⇒ Lots of parallelism

Downsides:

- Workload of  $2N$
- Communication!
- Workload stage dependent!

2022-02-16

Scan

└ Implementation

└ Up-Down Sweep

Up-Down Sweep

Dependency:

- Only between stages
- ⇒ Lots of parallelism

Downsides:

- Workload of  $2N$
- Communication!
- Workload stage dependent!

# Tiled Scan

Idea: Process input in independent chunks.

- Each chunk misses previous results  
⇒ Second pass over data.
- Workload:  $2N$

Our solution:

- Temporary vector for intermediate sums.
- Only one write to output.

2022-02-16

Scan  
└ Implementation  
└ Tiled Scan

Tiled Scan

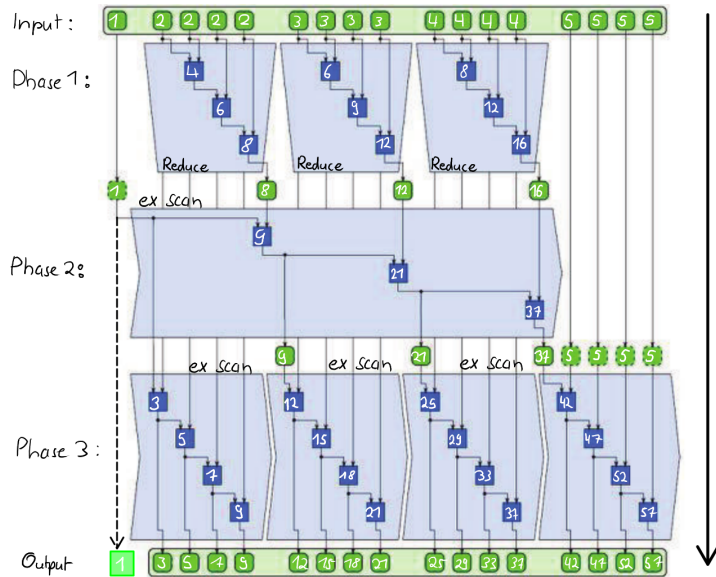
Idea: Process input in independent chunks.

- Each chunk misses previous results
  - ⇒ Second pass over data.
- Workload:  $2N$

Our solution:

- Temporary vector for intermediate sums.
- Only one write to output.

# Tiled Scan Schema




2022-02-16

Scan


└ Implementation

└ Tiled Scan Schema

Tiled Scan Schema



# Sequential Scan Results




2022-02-16

Scan

Implementation

Sequential Scan Results

Sequential Scan Results



# Segmented Scan

## Implementation

- Not present in STL!
- No reference implementations...

Solution: Wrapping the binary operation!

```
[binary_op](PairType left, PairType right){
    PairType new_right = right;
    if (not right.flag)
        new_right.value =
            binary_op(left.value, right.value);
    return new_right;
});
```

2022-02-16

Scan

└ Implementation

└ Segmented Scan

Segmented Scan  
Implementation

- Not present in STL!
- No reference implementations...

Solution: Wrapping the binary operation!

```
[binary_op](PairType left, PairType right){
    PairType new_right = right;
    if (not right.flag)
        new_right.value =
            binary_op(left.value, right.value);
    return new_right;
});
```

# Segmented Scan

## Solution

Works for:

- STL Scans
- Most inclusive scans

Challenge: Exclusive Scan

- Exclusive Segmented is complex
- Custom solution for each variant

2022-02-16

Scan  
└ Implementation  
└ Segmented Scan

Segmented Scan  
Solution


Works for:

- STL Scans
- Most inclusive scans

Challenge: Exclusive Scan

- Exclusive Segmented is complex
- Custom solution for each variant

# Sequential Segmented Scan Results




2022-02-16

Scan


Implementation

Sequential Segmented Scan Results





# Parallel Scan Results




2022-02-16


Scan  
└ Implementation

└ Parallel Scan Results

Parallel Scan Results



# Parallel Segmented Scan Results




2022-02-16

Scan

Implementation

Parallel Segmented Scan Results



# Intermediate Results

Ranking:

- 1 Library provided implementations
- 2 Tiled Scan
- 3 Up-Down Sweeping Scan

Remarks:

- OpenMP  $\geq$  TBB (if available)
- Up-Down Sweep is slow

Can we do better?

2022-02-16

Scan  
└ Optimizations

└ Intermediate Results

Ranking:  
1 Library provided implementations  
2 Tiled Scan  
3 Up-Down Sweeping Scan  
Remarks:  
• OpenMP  $\geq$  TBB (if available)  
• Up-Down Sweep is slow  
Can we do better?

# Algorithmic Optimization

Initial Goal: functional correctness.

Algorithmic Optimizations:

- Loop-Fusion:
  - Up-Down Sweep
  - Exclusive Segmented Scan
- Limiting Memory Accesses
- General clean up

Performance gain  $\sim$  1-5 GB/s!

2022-02-16

Scan  
└ Optimizations

└ Algorithmic Optimization

Initial Goal: functional correctness.

Algorithmic Optimizations:

- Loop-Fusion:
  - Up-Down Sweep
  - Exclusive Segmented Scan
- Limiting Memory Accesses
- General clean up

Performance gain  $\sim$  1-5 GB/s!

# Data Locality

Ensure that the data generated is local to the node:

```
std::vector<float> data(N);
#pragma omp parallel for schedule(static)
for (size_t i = 0; i < data.size(); i++)
{
    data[i] = rand();
}
```

- The performance gain by using data local structures is likely to be small due to the warmup of Catch2

2022-02-16

Scan  
└ Optimizations  
└ Data Locality

Data Locality

Ensure that the data generated is local to the node:

```
std::vector<float> data(N);
#pragma omp parallel for schedule(static)
for (size_t i = 0; i < data.size(); i++)
{
    data[i] = rand();
}
```

- The performance gain by using data local structures is likely to be small due to the warmup of Catch2

# OpenMP Scheduling




Fig. 1: Different Scheduling Strategies for 100 Iterations and 4 Threads

2022-02-16

Scan

└ Optimizations

└ OpenMP Scheduling

OpenMP Scheduling





Fig. 1: Different Scheduling Strategies for 100 Iterations and 4 Threads

- Static: Round-robin
- Dynamic: Constant Chunk Size
- Guided: Variable Chunk Size
- Auto: Compiler and/or Runtime System
- Cache Fusion if:
  1. single parallel region executes all of the maps to be fused.
  2. The loop for each map has the same bounds and chunk size.
  3. Each loop uses the static scheduling mode, either as implied by the environment or explicitly specified.
- Export or during runtime

# OpenMP Scheduling - Results MP-Media

mp-media OpenMP Provided Inclusive Scan with Different Scheduling (gnu)




2022-02-16

Scan

└ Optimizations


└ OpenMP Scheduling - Results MP-Media



- Static and Guided perform best
- Expected due to overhead

# OpenMP Scheduling - Results Ziti-Rome

ziti-rome OpenMP Provided Inclusive Scan with Different Scheduling (gnu)




2022-02-16

Scan

└ Optimizations

└ OpenMP Scheduling - Results Ziti-Rome

OpenMP Scheduling - Results Ziti-Rome



- More severe difference
- Static least overhead
- Static Bound size and Distribution of Work
- $\Rightarrow$  Static best



# TBB Partitioning

## TBB parallel constructs used:

- `parallel_scan`
- `parallel_for`

## available partitioners:

- `auto_partitioner`
- `affinity_partitioner`
- `simple_partitioner`
- `static_partitioner`

2022-02-16

Scan  
└ Optimizations

└ TBB Partitioning


**TBB parallel constructs used:**

- `parallel_scan`
- `parallel_for`

**available partitioners:**

- `auto_partitioner`
- `affinity_partitioner`
- `simple_partitioner`
- `static_partitioner`

## Performance (inclusive scan)




2022-02-16

Scan


└ Optimizations

└ Performance (inclusive scan)

Performance (inclusive scan)



## Performance (inclusive scan)




2022-02-16

Scan

└ Optimizations

└ Performance (inclusive scan)

Performance (inclusive scan)



# Vectorization

## Requirements:

- No loop carried dependency
- Loop bounds
- No jumps in code

## Realising it:

- `#pragma omp simd`
- Compiling with `-O3`
- Using Intel `lxc` Compiler

2022-02-16

Scan  
└─ Optimizations  
    └─ Vectorization

Vectorization


Requirements:

- No loop carried *dependency*
- Loop bounds
- No jumps in code

Realising it:

- `#pragma omp simd`
- Compiling with `-O3`
- Using Intel `lxc` Compiler

# Vectorization - Results MP-Media




2022-02-16

Scan

└ Optimizations


└ Vectorization - Results MP-Media

Vectorization - Results MP-Media



- Difference: Annotation of simd to loop
- Both compiled with -O3
- Possibly auto vectorization
- ICX much worse for MP-Media
- SIMD more stable

## Vectorization - Results Ziti-Rome




2022-02-16

Scan

└ Optimizations

└ Vectorization - Results Ziti-Rome



- Possibly auto vectorization
- ICX much better for ziti-rome
- SIMD better performance
- $\Rightarrow$  SIMD

# Summary

Library Provided Scans are fastest

2022-02-16

- Scan
  - Summary
    - Summary

Library Provided Scans are fastest

# Summary

Library Provided Scans are fastest

Optimization done:

- Algorithmic
- Data Locality
- Scheduler & Partitioner
- Vectorization

We have

- 4 versions of Scan
- 3 different algorithms
- 2 parallelization libraries + sequential

⇒ 36 Versions to keep track of!

2022-02-16

Scan  
└ Summary

└ Summary

Optimization done:

- Algorithmic
- Data Locality
- Scheduler & Partitioner
- Vectorization

We have

- 4 versions of Scan
- 3 different algorithms
- 2 parallelization libraries + sequential

⇒ 36 Versions to keep track of!