

Scan

Parallel Algorithm Design WS21/22

N. Kochendörfer, C. Alles, S. Proß

February 17, 2022

Table of Contents

- 1 Scan Theory
- 2 Implementation
- 3 Optimizations
- 4 Summary

2022-02-17

Scan

└─ Table of Contents

Table of Contents

- 1 Scan Theory
- 2 Implementation
- 3 Optimizations
- 4 Summary

Scan Theory

- Synonyms: prefix sum, cumulative sum or scan
- inclusive and exclusive version
- further specialization: segmented scan

2022-02-17

Scan
└ Scan Theory
 └ Scan Theory


Scan Theory

- Synonyms: prefix sum, cumulative sum or scan
- inclusive and exclusive version
- further specialization: segmented scan

Inclusive Scan

Prefix Sum
P of A

Store
Sum


<https://williamjrjibeiro.com/?p=132>


2022-02-17

Scan

└ Scan Theory

└ Inclusive Scan

Inclusive Scan



Inclusive vs. Exclusive scan

X

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 3 | 4 | 6 | 3 | 8 | 7 | 5 | 4 |
|---|---|---|---|---|---|---|---|

Y

| | | | | | | | |
|---|---|---|----|----|----|----|----|
| 0 | 3 | 7 | 13 | 16 | 24 | 31 | 36 |
|---|---|---|----|----|----|----|----|

Exclusive Scan

Y

| | | | | | | | |
|---|---|----|----|----|----|----|----|
| 3 | 7 | 13 | 16 | 24 | 31 | 36 | 40 |
|---|---|----|----|----|----|----|----|

Inclusive Scan

<https://livebook.manning.com/book/parallel-and-high-performance-computing/chapter-5/v-11/>

2022-02-17

Scan

└ Scan Theory

└ Inclusive vs. Exclusive scan

Inclusive vs. Exclusive scan

X

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 3 | 4 | 6 | 3 | 8 | 7 | 5 | 4 |
|---|---|---|---|---|---|---|---|

Y

| | | | | | | | |
|---|---|---|----|----|----|----|----|
| 0 | 3 | 7 | 13 | 16 | 24 | 31 | 36 |
|---|---|---|----|----|----|----|----|

Exclusive Scan

Y

| | | | | | | | |
|---|---|----|----|----|----|----|----|
| 3 | 7 | 13 | 16 | 24 | 31 | 36 | 40 |
|---|---|----|----|----|----|----|----|

Inclusive Scan

<https://livebook.manning.com/book/parallel-and-high-performance-computing/chapter-5/v-11/>

Segmented Variant

| | | | | | | |
|---|---|---|---|---|---|------------------|
| 1 | 2 | 3 | 4 | 5 | 6 | input |
| 1 | 0 | 0 | 1 | 0 | 1 | flag bits |
| 1 | 3 | 6 | 4 | 9 | 6 | segmented scan + |

https://en.wikipedia.org/wiki/Segmented_scan

2022-02-17

Scan

└ Scan Theory

└ Segmented Variant

| | | | | | | |
|---|---|---|---|---|---|------------------|
| 1 | 2 | 3 | 4 | 5 | 6 | input |
| 1 | 0 | 0 | 1 | 0 | 1 | flag bits |
| 1 | 3 | 6 | 4 | 9 | 6 | segmented scan + |

https://en.wikipedia.org/wiki/Segmented_scan

Implementation

STL Algorithm

STL provides:

- `std::inclusive_scan`
- `std::exclusive_scan`

Essentially equivalent to:

```
float sum = 0;
for (size_t i = 0; i < N; i++)
{
    sum += input[i];
    output[i] = sum;
}
```

⇒ Sequential to a fault!

2022-02-17

Scan

└ Implementation

└ Implementation

Implementation

STL Algorithm

STL provides:

- `std::inclusive_scan`
- `std::exclusive_scan`

Essentially equivalent to:

```
float sum = 0;
for (size_t i = 0; i < N; i++)
{
    sum += input[i];
    output[i] = sum;
}
```

⇒ Sequential to a fault!

Implementation

Alternatives

Alternatives to STL:

- OpenMP: scan pragma
- TBB: parallel_scan function

Alternative Algorithms:

- Up-Down Sweeping Scan
- Tiled Scan

2022-02-17

Scan

└ Implementation

└ Implementation

Implementation
Alternatives

Alternatives to STL:

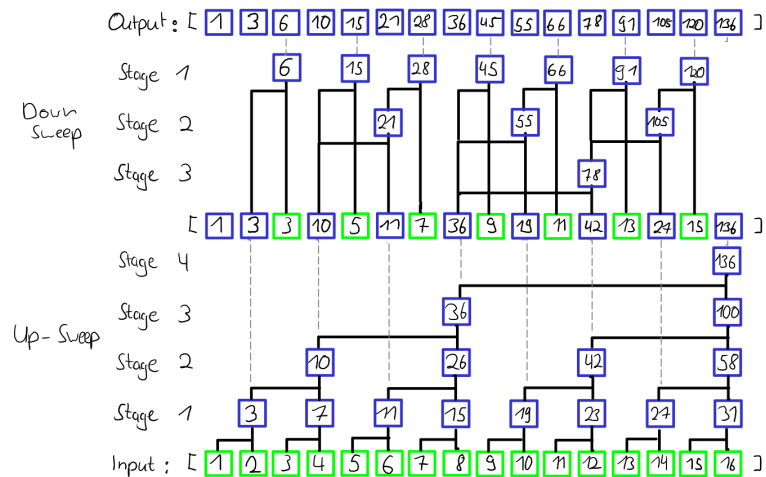
- OpenMP: scan pragma
- TBB: parallel_scan function

Alternative Algorithms:

- Up-Down Sweeping Scan
- Tiled Scan

Up-Down Sweep

Schema Inclusive




2022-02-17

Scan

Implementation

Up-Down Sweep

Up-Down Sweep
Schema Inclusive



Up-Down Sweep

Dependency:

- Only between stages

⇒ Lots of parallelism

Downsides:

- Workload of $2N$
- Communication!
- Workload stage dependent!

2022-02-17

Scan

└ Implementation

└ Up-Down Sweep

Up-Down Sweep

Dependency:

- Only between stages
- ⇒ Lots of parallelism

Downsides:

- Workload of $2N$
- Communication!
- Workload stage dependent!

Tiled Scan

Idea: Process input in independent chunks.

- Each chunk misses previous results
⇒ Second pass over data.
- Workload: $2N$

Our solution:

- Temporary vector for intermediate sums.
- Only one write to output.

2022-02-17

Scan
└ Implementation
└ Tiled Scan

Tiled Scan

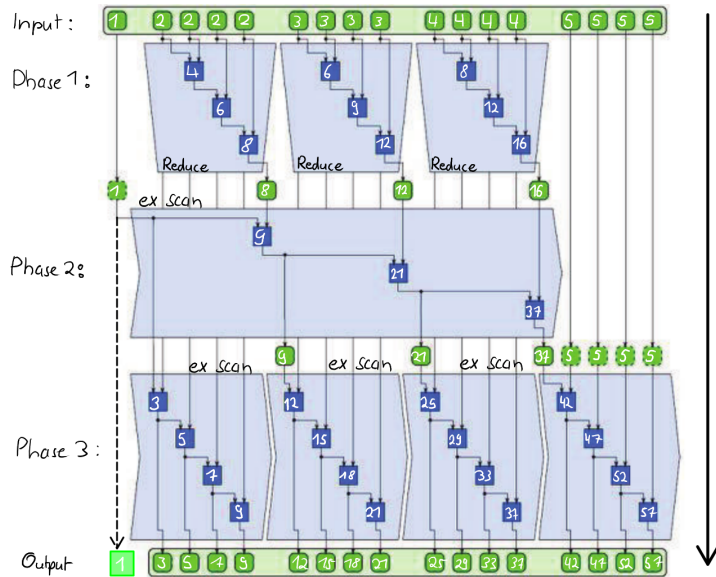
Idea: Process input in independent chunks.

- Each chunk misses previous results
 - ⇒ Second pass over data.
- Workload: $2N$

Our solution:


- Temporary vector for intermediate sums.
- Only one write to output.

Tiled Scan Schema




2022-02-17

Scan
└ Implementation
└ Tiled Scan Schema



Sequential Scan Results




2022-02-17

Scan

Implementation

Sequential Scan Results

Sequential Scan Results



Segmented Scan

Implementation

- Not present in STL!
- No reference implementations...

Solution: Wrapping the binary operation!

```
[binary_op](PairType left, PairType right){
    PairType new_right = right;
    if (not right.flag)
        new_right.value =
            binary_op(left.value, right.value);
    return new_right;
});
```

2022-02-17

Scan

└ Implementation

└ Segmented Scan

Segmented Scan
Implementation

- Not present in STL!
- No reference implementations...

Solution: Wrapping the binary operation!

```
[binary_op](PairType left, PairType right){
    PairType new_right = right;
    if (not right.flag)
        new_right.value =
            binary_op(left.value, right.value);
    return new_right;
});
```

Segmented Scan

Solution

Works for:

- STL Scans
- Most inclusive scans

Challenge: Exclusive Scan

- Exclusive Segmented is complex
- Custom solution for each variant

2022-02-17

Scan
└ Implementation
└ Segmented Scan

Segmented Scan
Solution


Works for:

- STL Scans
- Most inclusive scans

Challenge: Exclusive Scan

- Exclusive Segmented is complex
- Custom solution for each variant

Sequential Segmented Scan Results




2022-02-17

Scan


Implementation

Sequential Segmented Scan Results

Sequential Segmented Scan Results



Parallel Scan Results




2022-02-17

Scan


Implementation

Parallel Scan Results

Parallel Scan Results



Parallel Segmented Scan Results




2022-02-17

Scan

Implementation

Parallel Segmented Scan Results

Parallel Segmented Scan Results



Intermediate Results

Ranking:

- ① Library provided implementations
- ② Tiled Scan
- ③ Up-Down Sweeping Scan

Remarks:

- OpenMP \geq TBB (if available)
- Up-Down Sweep is slow

Can we do better?

2022-02-17

Scan
└ Optimizations

└ Intermediate Results

Intermediate Results

Ranking:
 ① Library provided implementations
 ② Tiled Scan
 ③ Up-Down Sweeping Scan
 Remarks:
 • OpenMP \geq TBB (if available)
 • Up-Down Sweep is slow
 Can we do better?

Algorithmic Optimization

Initial Goal: functional correctness.

Algorithmic Optimizations:

- Loop-Fusion:
 - Up-Down Sweep
 - Exclusive Segmented Scan
- Limiting Memory Accesses
- General clean up

Performance gain \sim 1-5 GB/s!

2022-02-17

Scan
└ Optimizations

└ Algorithmic Optimization

Initial Goal: functional correctness.

Algorithmic Optimizations:

- Loop-Fusion:
 - Up-Down Sweep
 - Exclusive Segmented Scan
- Limiting Memory Accesses
- General clean up

Performance gain \sim 1-5 GB/s!

Data Locality

Ensure that the data generated is local to the node:

```
std::vector<float> data(N);
#pragma omp parallel for schedule(static)
for (size_t i = 0; i < data.size(); i++)
{
    data[i] = rand();
}
```

- The performance gain by using data local structures is likely to be small due to the warmup of Catch2

2022-02-17

Scan
└ Optimizations
└ Data Locality

Data Locality

Ensure that the data generated is local to the node:

```
std::vector<float> data(N);
#pragma omp parallel for schedule(static)
for (size_t i = 0; i < data.size(); i++)
{
    data[i] = rand();
}
```

- The performance gain by using data local structures is likely to be small due to the warmup of Catch2

OpenMP Scheduling




Fig. 1: Different Scheduling Strategies for 100 Iterations and 4 Threads

2022-02-17

Scan

└ Optimizations

└ OpenMP Scheduling

OpenMP Scheduling





Fig. 1: Different Scheduling Strategies for 100 Iterations and 4 Threads

- Static: Round-robin
- Dynamic: Constant Chunk Size
- Guided: Variable Chunk Size
- Auto: Compiler and/or Runtime System
- Cache Fusion if:
- Export or during runtime

OpenMP Scheduling - Results MP-Media

mp-media OpenMP Provided Inclusive Scan with Different Scheduling (gnu)




2022-02-17

Scan

└ Optimizations


└ OpenMP Scheduling - Results MP-Media



- Static and Guided perform best
- Expected due to overhead

OpenMP Scheduling - Results Ziti-Rome

ziti-rome OpenMP Provided Inclusive Scan with Different Scheduling (gnu)




2022-02-17

Scan

└ Optimizations

└ OpenMP Scheduling - Results Ziti-Rome

OpenMP Scheduling - Results Ziti-Rome



- More severe difference
- Static least overhead
- Static Bound size and Distribution of Work
- \Rightarrow Static best

TBB Partitioning

TBB parallel constructs used:

- `parallel_scan`
- `parallel_for`

available partitioners:

- `auto_partitioner`
- `affinity_partitioner`
- `simple_partitioner`
- `static_partitioner`

2022-02-17

Scan
└ Optimizations

└ TBB Partitioning


TBB parallel constructs used:

- `parallel_scan`
- `parallel_for`

available partitioners:

- `auto_partitioner`
- `affinity_partitioner`
- `simple_partitioner`
- `static_partitioner`

Performance (inclusive scan)




2022-02-17

Scan


└ Optimizations

└ Performance (inclusive scan)

Performance (inclusive scan)



Performance (inclusive scan)




2022-02-17

Scan

└ Optimizations

└ Performance (inclusive scan)

Performance (inclusive scan)



Vectorization

Requirements:

- No loop carried dependency
- Loop bounds
- No jumps in code

Realising it:

- `#pragma omp simd`
- Compiling with `-O3`
- Using Intel `lxc` Compiler

2022-02-17

Scan
└─ Optimizations
 └─ Vectorization

Vectorization


Requirements:

- No loop carried *dependency*
- Loop bounds
- No jumps in code

Realising it:

- `#pragma omp simd`
- Compiling with `-O3`
- Using Intel `lxc` Compiler

Vectorization - Results MP-Media




2022-02-17

Scan

└ Optimizations


└ Vectorization - Results MP-Media

Vectorization - Results MP-Media



- Difference: Annotation of simd to loop
- Both compiled with -O3
- Possibly auto vectorization
- ICX much worse for MP-Media
- SIMD more stable

Vectorization - Results Ziti-Rome




2022-02-17

Scan

└ Optimizations

└ Vectorization - Results Ziti-Rome



- Possibly auto vectorization
- ICX much better for ziti-rome
- SIMD better performance
- \Rightarrow SIMD

Summary

Library Provided Scans are fastest

2022-02-17

Scan

└ Summary

└ Summary

Library Provided Scans are fastest

Summary

Library Provided Scans are fastest

Optimization done:

- Algorithmic
- Data Locality
- Scheduler & Partitioner
- Vectorization

We have

- 4 versions of Scan
- 3 different algorithms
- 2 parallelization libraries + sequential

⇒ 36 Versions to keep track of!

2022-02-17

Scan
└ Summary

└ Summary

Optimization done:


- Algorithmic
- Data Locality
- Scheduler & Partitioner
- Vectorization

We have

- 4 versions of Scan
- 3 different algorithms
- 2 parallelization libraries + sequential

⇒ 36 Versions to keep track of!

Roofline Mp-Media




2022-02-17


Scan
└ Summary

└ Roofline Mp-Media

Roofline Mp-Media



Roofline Ziti-Rome



2022-02-17

Scan
└ Summary

└Roofline Ziti-Rome

